

## Module 4 Quiz

( Mayur Brijwani )

### Q1. Notebook

We'll start with the same notebook we ended up with in homework 1. We cleaned it a little bit and kept only the scoring part. You can find the initial notebook [here](#).

Run this notebook for the March 2023 data.

What's the standard deviation of the predicted duration for this dataset?

- 1.24
- 6.24
- 12.28
- 18.28

**ANS - 6.24**

#### Q1. Notebook

We'll start with the same notebook we ended up with in homework 1. We cleaned it a little bit and kept only the scoring part. You can find the initial notebook [here](#).

Run this notebook for the March 2023 data.

What's the standard deviation of the predicted duration for this dataset?

- 1.24
- 6.24
- 12.28
- 18.28

**ANS - 6.24**

```
[12]: y_pred.std()
```

```
[12]: 6.247488852238703
```

## Q2. Preparing the output

Like in the course videos, we want to prepare the dataframe with the output.

First, let's create an artificial `ride_id` column:

```
df['ride_id'] = f'{year:04d}/{month:02d}_' + df.index.astype('str')
```

Next, write the ride id and the predictions to a dataframe with results.

Save it as parquet:

```
df_result.to_parquet( output_file, engine='pyarrow', compression=None, index=False )
```

What's the size of the output file?

- 36M
- 46M
- 56M
- 66M

**ANS - 66M**

```
[15]: df_result = pd.DataFrame()
      df_result['ride_id']=df['ride_id']
      df_result['predicted_duration'] = y_pred

[16]: df_result.head()

[16]:
```

	ride_id	predicted_duration
0	2023/03_0	16.245906
1	2023/03_1	26.134796
2	2023/03_2	11.884264
3	2023/03_3	11.997720
4	2023/03_4	10.234486

```


[17]: df_result.to_parquet( output_file, engine='pyarrow', compression=None, index=False )

[19]: !ls -a output/
      .  ..  yellow_tripdata_2023-03.parquet

[21]: !du -h ./output/yellow_tripdata_2023-03.parquet
      66M  ./output/yellow_tripdata_2023-03.parquet
```

### Q3. Creating the scoring script

Now let's turn the notebook into a script.

Which command you need to execute for that?

**ANS - jupyter nbconvert --to script starter.ipynb**

```
Q3. Creating the scoring script

Now let's turn the notebook into a script.

Which command you need to execute for that?

ANS - jupyter nbconvert --to script starter.ipynb

[24]: !jupyter nbconvert --to script starter.ipynb

[NbConvertApp] Converting notebook starter.ipynb to script
[NbConvertApp] Writing 3469 bytes to starter.py

[25]: !ls

model.bin  output  starter.ipynb  starter.py
```

### Q4. Virtual environment

Now let's put everything into a virtual environment. We'll use pipenv for that.

Install all the required libraries. Pay attention to the Scikit-Learn version: it should be the same as in the starter notebook.

After installing the libraries, pipenv creates two files: `Pipfile` and `Pipfile.lock`. The `Pipfile.lock` file keeps the hashes of the dependencies we use for the virtual env.

What's the first hash for the Scikit-Learn dependency?

**ANS -**  
**sha256:057b991ac64b3e75c9c04b5f9395eaf19a6179244c089afd**  
**eabad98264bff37c**

- Installing pipenv environment

```
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$ pipenv install scikit-learn==1.5.0 pandas --python=3.12.3
Installing scikit-learn==1.5.0...
Resolving scikit-learn==1.5.0...
Added scikit-learn to Pipfile's [packages] ...
✓ Installation Succeeded
Installing pandas...
Resolving pandas...
Added pandas to Pipfile's [packages] ...
✓ Installation Succeeded
Pipfile.lock (51d3e4) out of date: run `pipfile lock` to update to (404c93)...
Running $ pipenv lock then $ pipenv sync.
Locking [packages] dependencies...
Building requirements...
Resolving dependencies...
✓ Success!
Locking [dev-packages] dependencies...
Updated Pipfile.lock (c7db7b4b9e01fe02b3009be3d27b5e4985f82c60a155562c3444ee85d8404c93)!
Installing dependencies from Pipfile.lock (404c93)...
All dependencies are now up-to-date!
Installing dependencies from Pipfile.lock (404c93)...
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$ pipenv install pyarrow fastparquet
```

```
shell
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$ pipenv install pyarrow fastparquet
Installing pyarrow...
Resolving pyarrow...
Added pyarrow to Pipfile's [packages] ...
✓ Installation Succeeded
Installing fastparquet...
Resolving fastparquet...
Added fastparquet to Pipfile's [packages] ...
✓ Installation Succeeded
Pipfile.lock (404c93) out of date: run `pipfile lock` to update to (c2a809)...
Running $ pipenv lock then $ pipenv sync.
Locking [packages] dependencies...
Building requirements...
Resolving dependencies...
✓ Success!
Locking [dev-packages] dependencies...
Updated Pipfile.lock (d8b66c10ebf6ba3e2d81a8bae54127901c7f9fbfb488cc99a71ef7e34dc2a809)!
Installing dependencies from Pipfile.lock (c2a809)...
All dependencies are now up-to-date!
Installing dependencies from Pipfile.lock (c2a809)...
```

Ln 10, Col 12 (11 select)

```
Dockerfile predict_without_pipeline.py score.ipynb score.py starter.ipynb Pipfile.lock predict_with_pipeline.py
module-4 > homework > {} Pipfile.lock > {} default > {} scikit-learn > [ ] hashes
18     "default": {
114         "python-dateutil": {
121             },
122         "pytz": {
123             "hashes": [
124                 "sha256:2a29735ea9c18baf14b448846bde5a48030ed267578472d8955cd0e7443a9812",
125                 "sha256:328171f4e3623139da4983451950b28e95ac706e13f3f2630a879749e7a8b319"
126             ],
127             "version": "==2024.1"
128         },
129         "scikit-learn": {
130             "hashes": [
131                 "sha256:057b991ac64b3e75c9c04b5f9395eaf19a6179244c089afdebaad98264bfb37c",
132                 "sha256:118a8d229a41158c9f90093e46b3737120a165181a1b58c03461447aa4657415",
133                 "sha256:12e40ac48555e6b551f0a0a5743cc94cc5a765c9513fe708e01f0aa001da2801",
134                 "sha256:174beb56e3e881c90424e21f576fa69c4ffcf5174632a79ab4461c4c960315ac",
135                 "sha256:1b94d6440603752b27842eda97f6395f570941857456c606eb1d638efdb38184",
136                 "sha256:1f77547165c0062551e5c250cefa3f03f2fc92c5e18668abd90bfc4be2e0bfff",
137                 "sha256:261fe334ca48f09ed64b8fae13f9b46cc43ac5f580c4a605cbb0a517456c8f71",
138                 "sha256:2a65af2d8a6cce4e163a7951a4cfbfa7fceb2d5c013a4b593686c7f16445cf9d",
```

## Q5. Parametrize the script

Let's now make the script configurable via CLI. We'll create two parameters: year and month.

Run the script for April 2023.

What's the mean predicted duration?

- 7.29
- 14.29
- 21.29
- 28.29

Hint: just add a print statement to your script.

**ANS: 14.29**

starter.py × requirements.txt \$ start.sh Dockerfile .../mlops model.bin ✓ \$ .env.aws

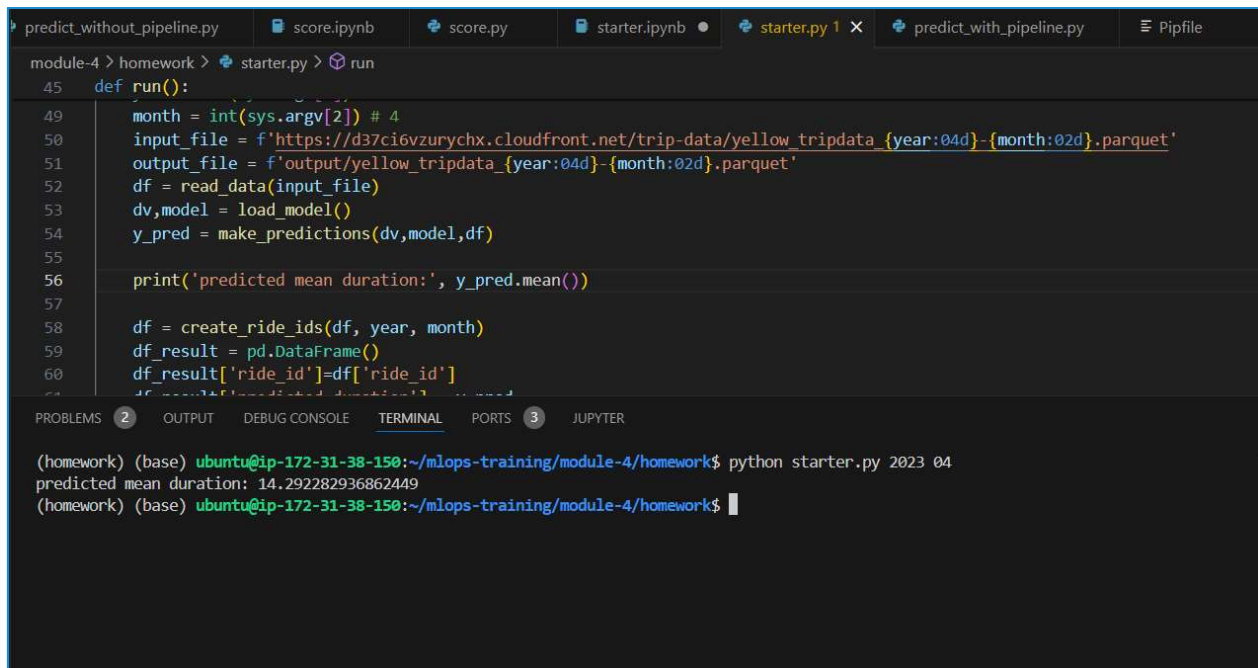
module-4 > homework > starter.py > create Ride IDs

```
1  #!/usr/bin/env python
2  # coding: utf-8
3
4  import pickle
5  import pandas as pd
6  import sys
7  import os
8  categorical = ['PULocationID', 'DOLocationID']
9
10
11 def load_model():
12     with open('../model.bin', 'rb') as f_in:
13         dv, model = pickle.load(f_in)
14     return dv, model
15
16
17 def read_data(filename):
18
19     df = pd.read_parquet(filename)
20     df['duration'] = df.tpep_dropoff_datetime - df.tpep_pickup_datetime
21     df['duration'] = df.duration.dt.total_seconds() / 60
22     df = df[(df.duration >= 1) & (df.duration <= 60)].copy()
23     df[categorical] = df[categorical].fillna(-1).astype('int').astype('str')
24     return df
25
26 def create Ride IDs(df, year, month):
27
28     df['ride_id'] = f'{year:04d}/{month:02d}_' + df.index.astype('str')
29     return df
30
31 def make_predictions(dv, model, df):
32
33     dicts = df[categorical].to_dict(orient='records')
34     X_val = dv.transform(dicts)
35     y_pred = model.predict(X_val)
36     return y_pred
37
```

```
starter.py × requirements.txt start.sh Dockerfile .../mlops model.bin ✓ .env.aws Dockerfile .../homework
module-4 > homework > starter.py > create Ride IDs

37
38
39
40 def run():
41
42     year = int(sys.argv[1]) # 2023
43     month = int(sys.argv[2]) # 4
44     input_file = f'https://d37ci6vzurychx.cloudfront.net/trip-data/yellow_tripdata_{year:04d}-{month:02d}.parquet'
45     #output_file = f'output/yellow_tripdata_{year:04d}-{month:02d}.parquet' //use this for storing output locally
46     output_file = f's3://module-04-output-mayur/yellow_tripdata_{year:04d}-{month:02d}.parquet'
47
48     df = read_data(input_file)
49     dv,model = load_model()
50     y_pred = make_predictions(dv,model,df)
51
52     print('predicted mean duration:', y_pred.mean())
53     print(y_pred)
54
55     df = create_ride_ids(df, year, month)
56     df_result = pd.DataFrame()
57     df_result['ride_id']=df['ride_id']
58     df_result['predicted_duration'] = y_pred
59
60     #os.makedirs('output', exist_ok=True)
61     df_result.to_parquet( output_file, engine='pyarrow', compression=None, index=False )
62
63 if __name__ == '__main__':
64     run()
65
66
67
68
69
70
```





The screenshot shows a JupyterLab environment with several tabs at the top: 'predict\_without\_pipeline.py', 'score.ipynb', 'score.py', 'starter.ipynb', 'starter.py' (active), and 'predict\_with\_pipeline.py'. The active tab 'starter.py' contains a Python function 'def run():'. The function takes a year and month as command-line arguments, reads a Parquet file from a CloudFront URL, loads a model, makes predictions, and prints the predicted mean duration. Below the code editor, the 'TERMINAL' tab is active, showing the command 'python starter.py 2023 04' and its output: 'predicted mean duration: 14.292282936862449'.

```
module-4 > homework > starter.py > run
45 def run():
49     month = int(sys.argv[2]) # 4
50     input_file = f'https://d37ci6vzurychx.cloudfront.net/trip-data/yellow_tripdata_{year:04d}-{month:02d}.parquet'
51     output_file = f'output/yellow_tripdata_{year:04d}-{month:02d}.parquet'
52     df = read_data(input_file)
53     dv,model = load_model()
54     y_pred = make_predictions(dv,model,df)
55
56     print('predicted mean duration:', y_pred.mean())
57
58     df = create_ride_ids(df, year, month)
59     df_result = pd.DataFrame()
60     df_result['ride_id']=df['ride_id']
61     df_result.to_parquet(output_file)
```

```
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$ python starter.py 2023 04
predicted mean duration: 14.292282936862449
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$
```

## Q6. Docker container

Finally, we'll package the script in the docker container. For that, you'll need to use a base image that we prepared

Now run the script with docker. What's the mean predicted duration for May 2023?

- 0.19
- 7.24
- 14.24
- 21.19

**ANS: 0.19**



```
service predict_without_pipeline.py score.ipynb score.py starter.ipynb starter.py Dockerfile ../hom
module-4 > homework > Dockerfile
1 FROM agrigorev/zoomcamp-model:mlops-2024-3.10.13-slim
2
3 RUN pip install -U pip
4 RUN pip install pipenv
5
6 WORKDIR /app
7
8 COPY [ "Pipfile", "Pipfile.lock", "./" ]
9
10 RUN pipenv install --system --deploy
11
12 COPY [ "starter.py", "starter.py" ]
13
14
15 ENTRYPOINT [ "python", "starter.py" ]

PROBLEMS 1 OUTPUT DEBUG CONSOLE TERMINAL PORTS 3 JUPYTER

Build an image from a Dockerfile
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$ docker run module_4_homework 2023 05
predicted mean duration: 0.19174419265916945
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$
(homework) (base) ubuntu@ip-172-31-38-150:~/mlops-training/module-4/homework$
```

## Bonus: upload the result to the cloud (Not graded)

Just printing the mean duration inside the docker image doesn't seem very practical. Typically, after creating the output file, we upload it to the cloud storage.

Modify your code to upload the parquet file to S3/GCS/etc.

```
starter.py requirements.txt start.sh Dockerfile .../ml_ops model.bin / .env.aws Doc
module-4 > homework > starter.py > run
46 def run():
47     output_file = f's3://module-04-output-mayur/yellow_tripdata_{year:04d}-{month:02d}.parquet'
48
49     df = read_data(input_file)
50     dv,model = load_model()
51     y_pred = make_predictions(dv,model,df)
52
53     print('predicted mean duration:', y_pred.mean())
54     print(y_pred)
55
56     df = create_ride_ids(df, year, month)
57     df_result = pd.DataFrame()
58     df_result['ride_id']=df['ride_id']
59     df_result['predicted_duration'] = y_pred
60
61     #os.makedirs('output', exist_ok=True)
62
63     df_result.to_parquet( output_file, engine='pyarrow', compression=None, index=False )
64
65 if __name__ == '__main__':
```

aws Services Search [Alt+S] N. Virginia

Amazon S3 > Buckets > module-04-output-mayur

### module-04-output-mayur Info

Objects Properties Permissions Metrics Management Access Points

**Objects (2) Info** [Refresh](#) [Copy S3 URI](#) [Copy URL](#) [Download](#) [Open](#) [Delete](#) [Actions](#) [Create folder](#) [Upload](#)

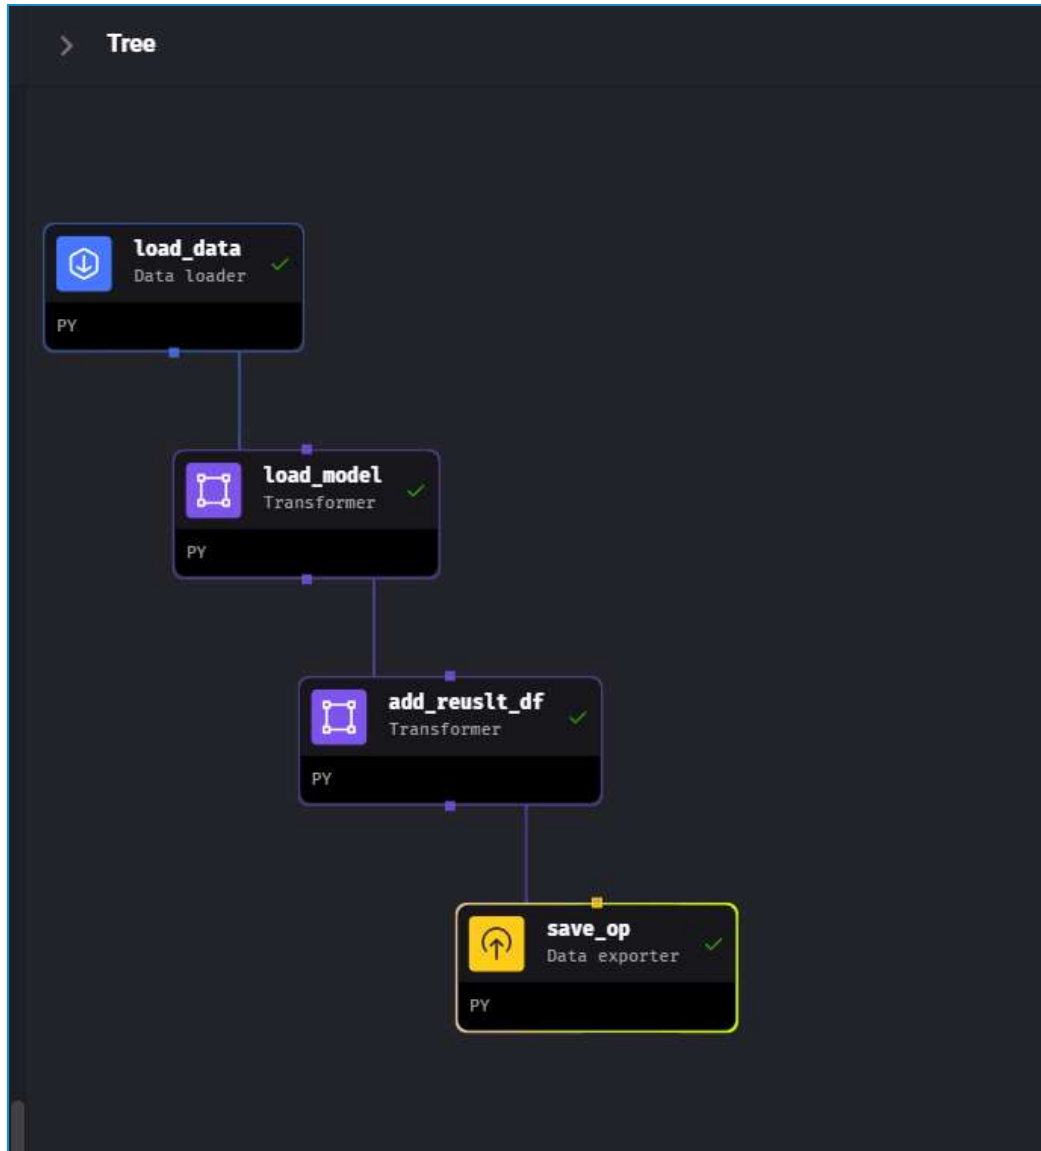
Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	<a href="#">yellow_tripdata_2023-04.parquet</a>	parquet	July 4, 2024, 20:00:34 (UTC+05:30)	63.1 MB	Standard
<input type="checkbox"/>	<a href="#">yellow_tripdata_2023-05.parquet</a>	parquet	July 4, 2024, 20:16:24 (UTC+05:30)	67.1 MB	Standard

## Bonus: Use Mage for batch inference

Here we didn't use any orchestration. In practice we usually do.

- Split the code into logical code blocks
- Use Mage to orchestrate the execution



- Creating data loader block to load data and add duration column

```
PY DATA LOADER load_data ← Edit parents

if 'data_loader' not in globals():
    from mage_ai.data_preparation.decorators import data_loader
if 'test' not in globals():
    from mage_ai.data_preparation.decorators import test

import pandas as pd

@data_loader
def load_data(*args, **kwargs):
    """
    Template code for loading data from any source.

    Returns:
    | Anything (e.g. data frame, dictionary, array, int, str, etc.)
    """
    # Specify your data loading logic here

    input_file = f"https://d37ci6vzurychx.cloudfront.net/trip-data/yellow_tripdata_{kwargs['year']}-{kwargs['month']}.parquet"
    df = pd.read_parquet(input_file)
    categorical = ['PULocationID', 'DOLocationID']
    df['duration'] = df.tpep_dropoff_datetime - df.tpep_pickup_datetime
    df['duration'] = df.duration.dt.total_seconds() / 60

    df = df[(df.duration ≥ 1) & (df.duration ≤ 60)].copy()

    df[categorical] = df[categorical].fillna(-1).astype('int').astype('str')

    return df

@test
def test_output(output, *args) → None:
    """
    Template code for testing the output of the block.
    """
    assert output is not None, 'The output is undefined'
```

1/1 tests passed.

OUTPUT 0

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID
0	1	2023-04-01T00:14:49.000	2023-04-01T00:45:01.000	2	4.9	1
1	2	2023-04-01T00:00:24.000	2023-04-01T00:56:19.000	1	21.89	2
2	1	2023-04-01T00:03:50.000	2023-04-01T00:14:42.000	2	1.3	1
3	1	2023-04-01T00:53:18.000	2023-04-01T01:01:28.000	1	1.5	1
4	2	2023-04-01T00:07:00.000	2023-04-01T00:17:16.000	2	1.49	1
5	1	2023-04-01T00:08:59.000	2023-04-01T00:15:39.000	6	1.2	1
6	2	2023-04-01T00:27:52.000	2023-04-01T00:43:07.000	1	8.61	1
7	2	2023-04-01T00:48:38.000	2023-04-01T01:08:37.000	1	3.88	1
8	1	2023-04-	2023-04-	0	8	1

3199715 rows x 20 columns

- Adding transformation block to load the vectorizer and model, transforming the validation set and creating predictions.

```
PY TRANSFORMER load_model ← 1 parent
@transformer
def transform(data):
    data → load_data

    if 'transformer' not in globals():
        from mage_ai.data_preparation.decorators import transformer
    if 'test' not in globals():
        from mage_ai.data_preparation.decorators import test
    import pickle
    import pandas as pd
    import sys
    import os

    @transformer
    def transform(df, *args, **kwargs):
        """
        Template code for a transformer block.

        Add more parameters to this function if this block has multiple parent blocks.
        There should be one parameter for each output variable from each parent block.

        Args:
            data: The output from the upstream parent block
            args: The output from any additional upstream blocks (if applicable)

        Returns:
            Anything (e.g. data frame, dictionary, array, int, str, etc.)
        """
        # Specify your transformation logic here
        with open('model.bin', 'rb') as f_in:
            dv, model = pickle.load(f_in)
        categorical = ['PULocationID', 'DOLocationID']

        dicts = df[categorical].to_dict(orient='records')
        X_val = dv.transform(dicts)
        y_pred = model.predict(X_val)
        print('predicted mean duration:', y_pred.mean())
        print(y_pred)
        return df, y_pred.tolist()
```

PY

TRANSFORMER

load\_model

← 1 parent

OUTPUT 0

OUTPUT 1

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID
0	1	2023-04-01T00:14:49.000	2023-04-01T00:45:01.000	2	4.9	1
1	2	2023-04-01T00:00:24.000	2023-04-01T00:56:19.000	1	21.89	2
2	1	2023-04-01T00:03:50.000	2023-04-01T00:14:42.000	2	1.3	1
3	1	2023-04-01T00:53:18.000	2023-04-01T01:01:28.000	1	1.5	1
4	2	2023-04-01T00:07:00.000	2023-04-01T00:17:16.000	2	1.49	1
5	1	2023-04-01T00:08:59.000	2023-04-01T00:15:39.000	6	1.2	1
6	2	2023-04-01T00:27:52.000	2023-04-01T00:43:07.000	1	8.61	1
7	2	2023-04-01T00:48:38.000	2023-04-01T01:08:37.000	1	3.88	1
8	1	2023-04-	2023-04-	0	8	1

3199715 rows x 20 columns

predicted mean duration: 14.292282936862449

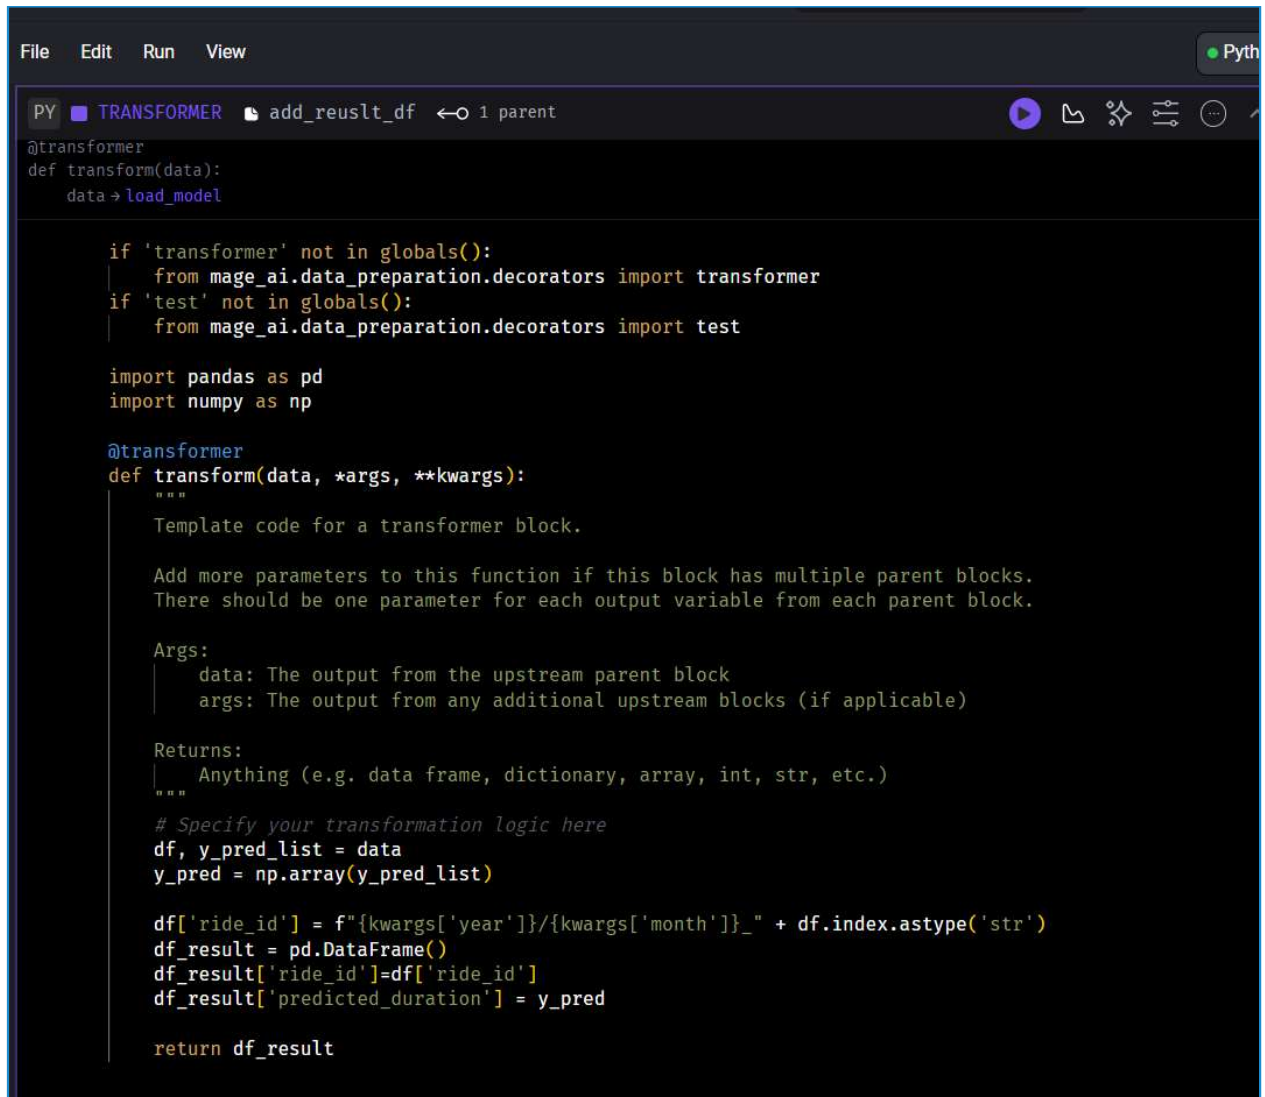
[16.13130203 32.26481598 12.25557264 ... 12.30118266 12.5560282 11.39399629]

58.971s

✓



- Adding df\_result dataframe to store predictions



```
File Edit Run View Python

PY TRANSFORMER add_result_df ← 1 parent

@transformer
def transform(data):
    data = load_model

    if 'transformer' not in globals():
        from mage_ai.data_preparation.decorators import transformer
    if 'test' not in globals():
        from mage_ai.data_preparation.decorators import test

    import pandas as pd
    import numpy as np

    @transformer
    def transform(data, *args, **kwargs):
        """
        Template code for a transformer block.

        Add more parameters to this function if this block has multiple parent blocks.
        There should be one parameter for each output variable from each parent block.

        Args:
            data: The output from the upstream parent block
            args: The output from any additional upstream blocks (if applicable)

        Returns:
            Anything (e.g. data frame, dictionary, array, int, str, etc.)
        """
        # Specify your transformation logic here
        df, y_pred_list = data
        y_pred = np.array(y_pred_list)

        df['ride_id'] = f"{kwargs['year']}/{kwargs['month']}_ " + df.index.astype('str')
        df_result = pd.DataFrame()
        df_result['ride_id'] = df['ride_id']
        df_result['predicted_duration'] = y_pred

        return df_result
```

PY TRANSFORMER add\_reuslt\_df ← 1 parent

```
"""
assert output is not None, 'The output is undefined'
```

1/1 tests passed.

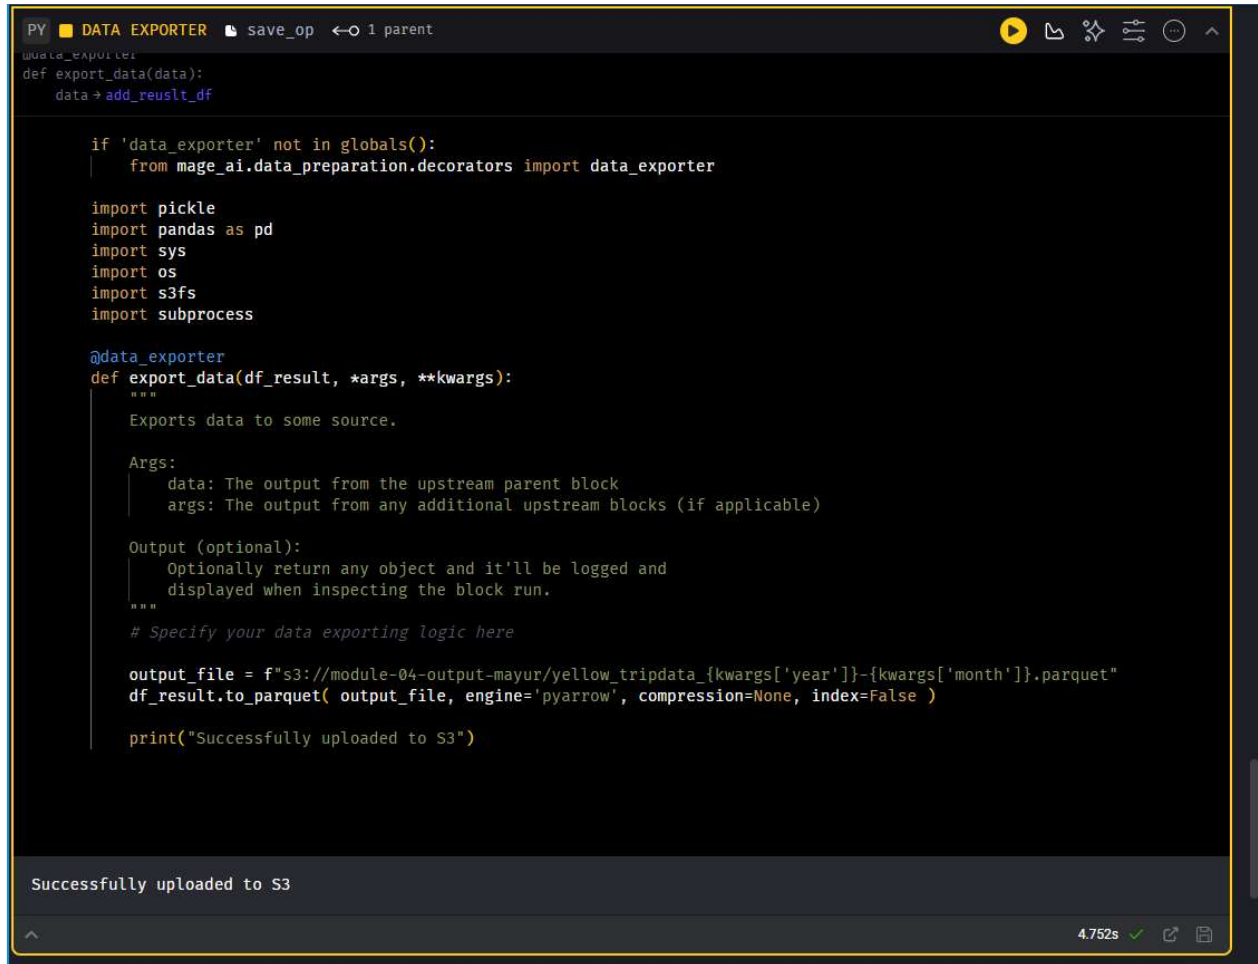
#### OUTPUT 0

	ride_id	predicted_duration
0	2023/04_0	16.1313020286
1	2023/04_1	32.2648159788
2	2023/04_2	12.255572639
3	2023/04_3	12.1336167818
4	2023/04_4	13.0719974054
5	2023/04_5	10.9383759277
6	2023/04_6	21.9629803406
7	2023/04_7	14.4678529806
8	2023/04_8	22.0833770565
9	2023/04_9	12.1992353141

3199715 rows x 2 columns

10.265s ✓

- Adding data exporter block to store results into S3 Bucket



The screenshot shows a Databricks notebook interface. At the top, the tab is labeled 'PY DATA EXPORTER' and the file name is 'save\_op'. Below the tab, the code is as follows:

```
data_exporter
def export_data(data):
    data = add_result_df

    if 'data_exporter' not in globals():
        from mage_ai.data_preparation.decorators import data_exporter

    import pickle
    import pandas as pd
    import sys
    import os
    import s3fs
    import subprocess

    @data_exporter
    def export_data(df_result, *args, **kwargs):
        """
        Exports data to some source.

        Args:
            data: The output from the upstream parent block
            args: The output from any additional upstream blocks (if applicable)

        Output (optional):
            Optionally return any object and it'll be logged and
            displayed when inspecting the block run.
        """
        # Specify your data exporting logic here

        output_file = f"s3://module-04-output-mayur/yellow_tripdata_{kwargs['year']}-{kwargs['month']}.parquet"
        df_result.to_parquet( output_file, engine='pyarrow', compression=None, index=False )

        print("Successfully uploaded to S3")
```

At the bottom of the notebook, a status bar shows 'Successfully uploaded to S3' and a timer of '4.752s'.

## module-04-output-mayur Info

**Objects** | Properties | Permissions | Metrics | Management | Access Points

Objects (5) Info

Copy S3 URI

Copy URL

Download

Open

Delete

Actions






Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

< 1 >

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	 <a href="#">yellow_tripdata_2023-04.parquet</a>	parquet	July 5, 2024, 16:46:47 (UTC+05:30)	63.1 MB	Standard
<input type="checkbox"/>	 <a href="#">yellow_tripdata_2023-05.parquet</a>	parquet	July 4, 2024, 20:16:24 (UTC+05:30)	67.1 MB	Standard
<input type="checkbox"/>	 <a href="#">yellow_tripdata_2023-06.parquet</a>	parquet	July 5, 2024, 13:52:26 (UTC+05:30)	63.3 MB	Standard
<input type="checkbox"/>	 <a href="#">yellow_tripdata_2023-07.parquet</a>	parquet	July 5, 2024, 13:06:23 (UTC+05:30)	55.7 MB	Standard
<input type="checkbox"/>	 <a href="#">yellow_tripdata_2023-08.parquet</a>	parquet	July 5, 2024, 16:41:15 (UTC+05:30)	54.1 MB	Standard